

Wall256

SU(3) Yang-Mills Strong-Coupling Conditional Reduction

A Complete Technical Report

History -- Architecture -- Barrier -- Experiments -- Conjectures -- Audit

David J. Fox | ORCID 0009-0008-1290-6105 | davidjfox998@gmail.com

Opera Numerorum | June 2026 | Lean 4 / Mathlib v4.12.0

HONESTY NOTICE -- READ FIRST This report makes **NO** claim that the Yang-Mills mass gap is proven. YM_STATUS = OPEN. Clay Surface #1: OPEN. No $\mu > 0$. No spectral gap. SORRY = 0. AXIOMS = {propext, Classical.choice, Quot.sound}. The entire mathematical content rests on **THREE** named open hypotheses (H1, H2, H3). None is discharged. The conclusion holds **ONLY IF** all three hold. Commit: 8eeab54 | Towers/YM/Wall256_Scaffold.lean | NOT a brick

AUDIT FINDINGS (caught during pre-build certification -- June 2026) AUDIT-1: P5_genuine = 1000000001119 = 7 x 142857143017 -- COMPOSITE, NOT PRIME. TOWERS_YM_v2.3 claimed this was the '13-digit boundary prime'. It is not prime. Remarkable: 142857143017 begins with 142857, the cyclic decimal of 1/7 (the KP threshold). The boundary number is divisible by the threshold denominator 7. AUDIT-2: beta_0 bracket: our 7-moment partial sum cannot reproduce the CERT_Arb bracket; full 36-moment computation (performed in source workspace) is authoritative. CERT_Arb SHA: b5a9f0a7666a91f283a7d4531ae99dff2097c2cedef10424b77833b7bbc840d3

COMMIT = 8eeab54 | SORRY = 0 | YM_STATUS = OPEN

Contents

1. The Historical Problem: Yang-Mills Mass Gap
 2. The Architecture: Everything Reduces to One Inequality
 3. The Wall256 Theorem -- Verbatim Lean (commit 8eeab54)
 4. The Three Open Inputs: H_1 , H_2 , H_3
 5. Why the Bound Must Be Strict: $w_1 < 1/7$
 6. The β_0 Certificate -- $SU(3)$ Haar Integral
 7. The Bessel I Barrier -- Where mathlib Runs Out
 8. Surface #1 Installment A -- Peter-Weyl Infrastructure
 9. Surface #2 -- Polymer Expansion Scaffold
 10. The Z-Experiments -- Numerical Testing
 11. The 13-Module Witness-Collapse Audit
 12. The Conjecture Landscape -- Open Problems
 13. Audit Findings
 14. What IS and IS NOT Proven (Clay rigour checklist)
 15. Cryptographic Binding
-

1. The Historical Problem: Yang-Mills Mass Gap

The Yang-Mills existence and mass gap problem -- one of the seven Clay Millennium Prize Problems -- asks, for a compact simple Lie gauge group G (here $G = \text{SU}(3)$, the gauge group of QCD), for the construction of a quantum Yang-Mills theory on $\mathbb{R}^4$ satisfying the Wightman axioms, whose Hamiltonian has a STRICTLY POSITIVE spectral gap above the vacuum.

The mass gap $m > 0$ explains why the strong nuclear force is short-range despite being classically mediated by massless gluons -- the quantum theory acquires a mass through confinement. Proving this rigorously from first principles is the problem.

Three Equivalent Formulations

- (E1) Two-point exponential clustering:
 $|\langle W(C_1) W(C_2) \rangle - \langle W(C_1) \rangle \langle W(C_2) \rangle| \leq C * \exp(-\Delta * \text{dist}(C_1, C_2)), \quad \Delta > 0$
 $W(C)$ = Wilson loop (gauge-invariant trace of holonomy around C)
<. > = expectation in the Euclidean lattice YM measure
- (E2) Transfer-operator spectral gap:
 $\text{spec}(T_{\text{real}} | H_0) \text{ in } [0, \lambda_{1,1}], \quad \lambda_{1,1} < 1$
 $\Delta = -\log(\lambda_{1,1}) > 0$
 T_{real} = Wilson/heat-kernel transfer operator (one Euclidean time step)
- (E3) HasMassGap predicate:
 $\langle x, T x \rangle \geq (1-m) \|x\|^2 \text{ for all } x \text{ in } H_0, \quad m > 0$
Equivalently: $\lambda_{1,1} \leq \exp(-m)$

The HONEST GAP: a non-vacuous discharge of (E1)-(E2)-(E3) requires a REAL $\text{SU}(3)$ Wilson transfer operator with a verified $T_{\text{real}} > 0$. The ingredients -- $\text{SU}(3)$ character theory and a verified heat-kernel transfer operator with explicit spectral bounds -- are absent from mathlib v4.12.0. Every YM module in this repository that claimed to touch (E1)-(E3) was discharged at a degenerate witness (see Section 11).

2. The Architecture: Everything Reduces to One Inequality

The Osterwalder-Seiler strong-coupling analysis provides a rigorous lattice route to exponential clustering. The entire chain of implications collapses to a single condition on the single-site Haar weight:

THE HINGE INEQUALITY: $w_1 < 1/7$ $w_1 = \int_{\text{SU}(3)} \exp(-\beta * S) d(\text{Haar})$, $S = (3 - \text{Re tr } U) / 3$ If $w_1 < 1/7$, the polymer entropy series converges (KP criterion), and the lattice two-point function decays exponentially with rate $\Delta > 0$. This requires $\beta > \beta_0 = 2.079416880123\dots$ (CERT_Arb certified).

The Reduction Tower

$w_1 < 1/7$ single-site Haar weight is strictly small	[H1] OPEN
$\Rightarrow a(n) \leq \exp(-I * n), \quad I > \log(7) = 1.9459\dots$ Osterwalder-Seiler 1978 Thm 2.1: single-site smallness propagates to per-polymer activity rate $I > \log(7)$	[H2] OPEN
$\Rightarrow \sum_n N(n) * a(n) \leq \sum_n 7^n * \exp(-I * n)$ $\quad = \sum_n (7 * \exp(-I))^n < \text{infinity}$ Comparison test converges since $I > \log(7) \Leftrightarrow 7 * \exp(-I) < 1$ [kp_summable_of_truncatedActivity -- machine-checked]	PROVED
$\Rightarrow 0 < \rho < 1, \text{corr}(x, y) \leq C * \rho^{\text{sep}(x, y)}$ Brydges-Federbush: KP summability $\Rightarrow$ geometric two-point decay	[H3] OPEN
$\Rightarrow \Delta = -\log(\rho) > 0, \text{corr} \leq C * \exp(-\Delta * \text{sep})$ $\rho^d = \exp(-\Delta * d)$ algebra [mass_gap_pos_of_spectral_gap -- machine-checked]	PROVED
 CONDITIONAL CONCLUSION: abstract lattice two-point decay Holds ONLY IF H1 + H2 + H3 are ALL discharged. This is NOT a mass gap, NOT a Clay result, NOT a spectral gap.	

3. The Wall256 Theorem -- Verbatim Lean (commit 8eeab54)

```
-- Towers/YM/Wall256_Scaffold.lean (commit 8eeab54)
-- SORRY: 0 | AXIOMS: classical trio | NOT a brick | YM_STATUS: OPEN

theorem strong_coupling_decay_of_open_inputs
  {E : Type*} (corr sep : E -> E -> R) (C rho w1 : R)
  {N a : N -> R}
  (hN0 : forall n, 0 <= N n)
  (hN : forall n, N n <= (7 : R) ^ n)
  (hw1 : w1 < 1 / 7) -- H1 OPEN
  (hOS : w1 < 1 / 7 -> TruncatedActivityBound a) -- H2 OPEN
  (h_bridge : Summable (fun n : N => N n * a n) -> -- H3 OPEN
    0 < rho /\ rho < 1 /\
    forall x y, |corr x y| <= C * rho ^ (sep x y)) :
  exists Delta : R,
    0 < Delta /\ forall x y,
      |corr x y| <= C * Real.exp (-Delta * sep x y) :=
  su2_gap_of_truncatedActivity corr sep C rho hN0 hN (hOS hw1) h_bridge
```

Symbol	Meaning
E : Type*	Abstract type -- ANY space of sites. No real lattice built.
corr, sep	Abstract correlator and separation. No real Wilson loop.
N n <= 7^n	Polymer entropy: connected polymers of size n number at most 7^n (SU(3) coordination).
hw1 : w1 < 1/7	H1 OPEN -- single-site Haar weight strictly below the KP threshold.
hOS	H2 OPEN -- Osterwalder-Seiler: single-site smallness => activity rate I > log(7).
h_bridge	H3 OPEN -- Brydges-Federbush: KP summability => geometric two-point clustering.
TruncatedActivityBound a	(forall n, 0<=a n) AND (exists I, log(7)<I AND forall n, a n<=exp(-I)^n)
su2_gap_of_truncatedActivity	Already-proved combinator in Wall256_Note.lean. Group-agnostic. Asserts no gap.
Conclusion	EXISTS Delta>0 s.t. corr <=C*exp(-Delta*sep). CONDITIONAL on H1+H2+H3 only.

SORRY = 0. Proof body is one line: plug H1, H2, H3 into the proved combinator.

No real SU(3) operator, no real Wilson loop, no real lattice metric is built.

4. The Three Open Inputs: H1, H2, H3

H1 -- SU(3) single-site Haar weight (strict bound)

```
hw1 : w1 < 1/7
w1 = integral_{SU(3)} exp(-beta * S) dHaar, S = (3 - Re tr U)/3
```

The single-site Boltzmann weight $w1(\beta)$ decreases as β increases (Laplace transform of a non-degenerate non-negative random variable -- strictly monotone). The strict bound $w1 < 1/7$ holds for $\beta > \beta_0$, certified at β_0 in [2.079416880123, 2.079416880124] (CERT_Arb). Mathlib v4.12.0 cannot evaluate SU(3) Haar integrals, so this stays a named hypothesis on abstract $w1$.

STATUS: OPEN. Source: SU(3) character theory / verified cubature. Absent from mathlib v4.12.0.

H2 -- Osterwalder-Seiler 1978, Theorem 2.1 (Ursell / cluster step)

```
hOS : w1 < 1/7 -> TruncatedActivityBound a
```

Single-site smallness propagates, via the truncated (Ursell) cluster expansion, to a per-size connected-polymer activity bound: $a(n) \leq \exp(-I \cdot n)$ with $I > \log(7)$ STRICTLY. The cluster expansion is absent from mathlib. The strict rate $I > \log(7)$ (not $\geq$) is essential -- see Section 5.

STATUS: OPEN. Source: Osterwalder & Seiler, Ann. Phys. 110 (1978), Thm 2.1. Absent from mathlib v4.12.0. Hardest open leaf.

H3 -- Brydges-Federbush KP Bridge

h_bridge : Summable (N n * a n) ->
 $0 < \rho$ AND $\rho < 1$ AND forall x y, $|\text{corr } x \ y| \leq C * \rho^{\text{sep } x \ y}$

Kotecky-Preiss summability of the entropy-weighted polymer series implies geometric two-point clustering with spectral radius $\rho < 1$. Standard cluster-expansion theory (Friedli & Velenik, Stat. Mech. of Lattice Systems, Ch. 5), absent from mathlib v4.12.0. Carried as a hypothesis, NOT by sorry.

STATUS: OPEN. Source: Friedli & Velenik 2018, Ch. 5. Absent from mathlib v4.12.0.

5. Why the Bound Must Be Strict: $w_1 < 1/7$ (not $\leq$)

The strictness is not a formality -- it is the entire reason the polymer entropy series converges. The boundary case $w_1 = 1/7$ makes the comparison test fail (geometric series ratio = 1, series diverges).

Case	Rate	Geometric ratio	Series	Verdict
$w_1 < 1/7$	$I > \log(7)$	$7 \cdot \exp(-I) < 1$	$\text{sum } (7 \cdot \exp(-I))^n$	CONVERGES (KP ok)
$w_1 = 1/7$ (boundary, $\beta = 0.85$ excluded)	$I = \log(7)$	$7 \cdot \exp(-I) = 1$	$\text{sum } 1^n = \text{sum } 1$	DIVERGES (vacuous)
$w_1 > 1/7$	$I < \log(7)$	$7 \cdot \exp(-I) > 1$	grows without bound	DIVERGES

CERTIFIED: Convergence vs divergence at the boundary

```
log(7)          = 1.945910149055313
exp(-log(7)) = 0.142857142857143 = 1/7  [VERIFIED: identical to 15dp]

BOUNDARY (I = log 7): ratio = 7 * exp(-log7) = 1.000000000000000 => DIVERGES

I = log(7) + 0.001:  ratio = 0.999000  sum = 1000.5  [CONVERGES]
I = log(7) + 0.010:  ratio = 0.990050  sum = 100.5   [CONVERGES]
I = log(7) + 0.100:  ratio = 0.904837  sum = 10.5    [CONVERGES]
I = log(7) + 0.500:  ratio = 0.606531  sum = 2.5     [CONVERGES]

Heuristic at beta=2.07942 (just above beta_0):
w1 ~ 0.142853 < 1/7 = 0.142857 => I = -log(w1) ~ 1.945960 > log(7) = 1.945910
I - log(7) ~ 0.000050 (very narrow strict margin just above threshold)
```

Strict inequality $I > \log(7)$ required and verified at $\beta > \beta_0$.
Boundary $\beta = 0.85$ ($w_1 = 1/7$) is EXCLUDED -- the reduction is vacuous there.

6. The β_0 Certificate -- SU(3) Haar Integral

STATUS: VERIFIED_OUT_OF_TOWER (CERT_Arb, authoritative) NOT Lean. NOT trio-clean. Numerics only. Discharges ZERO Lean obligations. AUDIT-2 NOTE: Our 7-moment partial sum ($m_0..m_6$) cannot reproduce the CERT_Arb bracket. The full computation uses $N=36$ exact moments via constant-term extraction over the SU(3) torus. That computation was performed in the source workspace (TheoremaAureum143). We certify the qualitative results below; CERT_Arb is the authoritative enclosure.

The SU(3) Haar Integral Formula

```
w1(beta) = exp(-beta) * sum_{n=0}^{inf} (beta/3)^n * m_n / n!

m_n = <(Re tr U)^n>_{Haar on SU(3)} -- exact rational moments
computed by constant-term extraction over the SU(3) torus
(Weyl integration formula)

Rigorous tail bound: |R_N| <= beta^{N+1} / (N+1)! / (1 - beta/(N+2))
(from |Re tr U| <= 3 => |m_n| <= 3^n)

CERT_Arb: N=36 terms, tail <= 4.46e-32 (mpmath.iv dps=80, outward rounding)
```

Exact Moments m_n (from CERT_Arb -- machine-certified rationals)

n	m_n (exact rational)	m_n (decimal)
0	1	1.000000000000000
1	0	0.000000000000000 (SU(3) symmetry)
2	1/2	0.500000000000000
3	1/4	0.250000000000000
4	3/4	0.750000000000000

5	15/16	0.9375000000000000
6	65/32	2.0312500000000000

Certified w1 sign table (CERT_Arb, N=36, dps=80)

beta	w1(beta) enclosure	vs 1/7	Verdict
0.86	[0.432367..., 0.432367...]	> 1/7 = 0.142857...	D4 FAILS -- w1 >> 1/7
2.07941	[0.142857993..., 0.142857994...]	> 1/7	Above threshold
2.07942	[0.142856757..., 0.142856758...]	< 1/7	Below threshold
2.079416880123	certified > 1/7	> 1/7	Low bracket endpoint
2.079416880124	certified < 1/7	< 1/7	High bracket endpoint
beta_0 in [2.079416880123, 2.079416880124]		CROSSING CERTIFIED	LOCKED

D4 at beta=0.86: $w1(0.86) > 1/7$ CERTIFIED. D4 FAILS.

beta_0 in [2.079416880123, 2.079416880124]: CERTIFIED (CERT_Arb).

Monotonicity: $w1(\beta) = E[\exp(-\beta \cdot S)]$ is strictly decreasing (Laplace transform of non-degenerate $S \geq 0$). Crossing is unique.

This does NOT produce a Lean proof object. hw1 remains an open named hypothesis in Wall256_Scaffold.lean.

7. The Bessel I Barrier -- Where mathlib Runs Out

The Haar integral $w_1(\beta)$ -- the hinge of the entire Wall256 reduction -- is computed via the $SU(3)$ character expansion, which involves modified Bessel functions of the first kind. This is the point where the formal proof hits a genuine barrier: mathlib v4.12.0 has no $SU(3)$ Haar measure, no character theory, and no Bessel function library.

How Bessel Functions Enter

$SU(3)$ heat kernel at the identity:
 $K_\beta(I) = \sum_{\{m,n\} \in N^2} \dim(m,n)^2 * \exp(-\beta * C_2(m,n))$

$\dim(m,n) = (m+1)(n+1)(m+n+2)/2$ (Weyl dimension formula)
 $C_2(m,n) = m^2 + n^2 + mn + 3m + 3n$ (quadratic Casimir, un-normalized)

Character expansion of $\text{Re tr } U$ on $SU(2)$ analogy:
 $\langle \exp(i * \theta) \rangle_{SU(2)}$ involves $I_0(\beta)$ and $I_1(\beta)$
where $I_k(x) = \sum_{j=0}^{\infty} (x/2)^{2j+k} / (j! * (j+k)!)$

For $SU(3)$: moments $m_n = \langle (\text{Re tr } U)^n \rangle$ computed by constant-term extraction over the maximal torus via the Weyl integration formula.
The closed-form expressions involve $SU(3)$ Bessel/character integrals.

Certified Bessel I values (verified against Z_BESSEL_I_COMPLETE CSV)

$n \setminus x$	$x = 0.5$	$x = 1.0$	$x = 2.0$	$x = 5.0$	$x = 10.0$
I_0	1.063483370741	1.266065877752	2.279585302336	27.239871823604	2815.716628466
I_1	0.257894305391	0.565159103992	1.590636854637	24.335642142451	2670.988303701
I_2	0.031906149178	0.135747669767	0.688948447699	17.505614966624	2281.518967726
I_3	0.002645111969	0.022168424924	0.212739959240	10.331150169151	1758.380716611

ALL 20 Bessel I values verified against Z_BESSEL_I_COMPLETE CSV: all PASS (relative error < 1e-10). mpmath computation is deterministic.

What is absent from mathlib v4.12.0

Missing component	Status	Impact
$SU(3)$ Haar measure	ABSENT	Cannot state hw_1 as a Lean term
$SU(3)$ character theory	ABSENT	Cannot compute m_n in Lean
Modified Bessel functions I_k	ABSENT	No Bessel-based bound possible
Weyl integration formula	ABSENT	No torus decomposition of Haar
$w_1(\beta) < 1/7$ proof	ABSENT	H_1 / hw_1 stays open hypothesis

The Bessel I barrier is the primary reason H_1 (hw_1) cannot be discharged in Lean with the current mathlib. All numerical evidence is consistent with the claimed threshold; the formal proof gap is analytic, not arithmetic.

8. Surface #1 Installment A -- Peter-Weyl Infrastructure

While the full Haar integral remains outside mathlib, the Peter-Weyl heat-kernel series HAS been formalized. Two lemmas provide the summability infrastructure needed by any cluster-expansion attack on the mass gap.

Formalized Results (8 bricks, tasks #154 and #155)

Theorem	Lean name	Statement	Status
3.1	PeterWeyl_Summable_SU3	For every $t > 0$: $\dim(m,n)^2 * \exp(-t * C_2(m,n))$ is Summable (PROVED) in N^2 .	PROVED
3.2	Weyl_sum_SU3_le_envelope	For every $t > 0$, N in N : $S_N(t) \leq K(t)$ (truncation $\leq$ full tsb) (PROVED)	PROVED
3.3	Heat_kernel_envelope_ge_one	For every $t > 0$: $K(t) \geq 1$ (trivial rep (0,0) contributes 1). (PROVED)	PROVED

Verified $S_N(t)$ convergence (computed, not proved -- confirms formulas)


```

S_N(t=1), N= 5: sum = 1.340902 (21 pairs)
S_N(t=1), N=10: sum = 1.340902 (66 pairs) -- converged at N=5
S_N(t=1), N=20: sum = 1.340902 (231 pairs) -- stable to 6 dp
S_N(t=1), N=50: sum = 1.340902 (1326 pairs) -- no further contribution

dim(0,0)=1, C2(0,0)=0: PASS dim(1,0)=3, C2(1,0)=4: PASS
dim(0,1)=3 (anti-fundamental): PASS

```

Open gaps

Task	Description	Blocks
#156 OPEN	Varadhan small-t asymptotic: $K(t) \leq C t^{-4} \exp(-c/t)$	Per-plaquette activity bound, all of Surface #2
#157 OPEN	Tighter Casimir: $C_2(m,n) \geq c(m+n)^2$ as Lean lemma	Tech-debt; shortens Surface #1 proofs only

460 trio-clean bricks total. 8 bricks in [PeterWeyl.lean](#) + [PeterWeylHeat.lean](#).

Genesis seal: [eecbcd9a540aa7a2c90edd23827c73e4d1bb5af641d352f70a5de849b21f875f](#)

9. Surface #2 -- Polymer Expansion Scaffold

Surface #2 describes the polymer-expansion route to a convergent cluster expansion for 4D SU(3) lattice YM at small coupling. It is the intended consumer of the Peter-Weyl infrastructure from Surface #1. Every central inequality is OPEN.

```
STEP 1: Per-plaquette exponential activity [OPEN -- Lean sorry]
exists_c_per_plaquette_pw:
plaquette_activity_pw(beta,N,p) <= C * beta^{-4} * exp(-c/beta)
[blocked by Varadhan asymptotic, Task #156]

STEP 2: Kotecky-Preiss criterion [OPEN -- ~60 sorries in ClusterExpansion.lean]
kotecky_preiss_criterion:
exists mu s.t. sum_{gamma:'gamma'capta} w(gamma')*exp(mu(gamma')) <= mu(gamma)

STEP 3: Cluster expansion of log Z_Lambda [OPEN -- not formalized]
log Z_Lambda(beta) = sum_X Phi_T(X) * prod_{gamma in X} w(gamma)

STEP 4: Area-law Wilson loop decay [OPEN -- Wilson loops not defined]
|<W_C>| <= K * exp(-m * Area(C))

CONTINUUM LIMIT a -> 0 [SEPARATE SURFACE -- not attempted]
```

YM tower Status: OPEN. No step above is formalized. ~60 sorries remain in Towers/Attempts/ClusterExpansion.lean.

10. The Z-Experiments -- Numerical Testing

The Z-experiment harness tests two computation modes: T=1 (tool-assisted, deterministic) and T=0 (LLM-direct prediction). The harness measures reliability, not mathematical content. All T=0 Bessel runs are NOT_RUN_NO_API in this batch.

Z_POLYMER: LLM-direct vs tool-assisted on bit-string polymers

Input	L	Zero-run	T=1 errors	T=0 errors
10111111	8	1	0/100 (0.00)	78/100 (0.78)
10000111	8	4	0/100 (0.00)	100/100 (1.00)
1001111111111111	16	2	0/100 (0.00)	52/100 (0.52)
1000000001111111	16	8	0/100 (0.00)	100/100 (1.00)
100011...11 (24b)	24	3	0/100 (0.00)	100/100 (1.00)
100000...11 (24b)	24	12	0/100 (0.00)	100/100 (1.00)
1000001...11 (40b)	40	5	0/100 (0.00)	100/100 (1.00)
1000000...11 (40b)	40	20	0/100 (0.00)	100/100 (1.00)

Z_BOSTCONNES: LLM-direct vs tool-assisted on KMS state

beta_KMS	T=1 errors/100	T=0 errors/100
1.5	0 (0.00)	100 (1.00)
2.0	0 (0.00)	76 (0.76)
2.5	0 (0.00)	24 (0.24)
3.0	0 (0.00)	100 (1.00)
5.0	0 (0.00)	73 (0.73)
8.0	0 (0.00)	30 (0.30)

Experiment	T=1 result	T=0 result	Finding
Z_BESSEL_I (20 cells)	0/400 errors	NOT_RUN (NO-API)	Bessel I: mpmath is exact and deterministic
Z_POLYMER (8 inputs)	0/800 errors	52-100% errors	LLM unreliable for long zero-run bit strings
Z_BOSTCONNES (8 cells)	0/800 errors	24-100% errors	LLM erratic for KMS evaluation (no pattern)

Z_POLYMER_RESULTS (101)	0 errors	NOT_RUN	Tool-assisted reproduction fully deterministic
-------------------------	----------	---------	--

T=1 (tool-assisted): 0 errors across ALL experiments. Deterministic.

T=0 (LLM-direct): 24-100% error rates. Error rate increases with zero-run length and bit count. Not suitable for formal certification.

11. The 13-Module Witness-Collapse Audit (PDF 263, 2026-06-03)

On 2026-06-03 a repair audit converted all 13 YM/OS modules from 'proved theorems' to named open Props and de-registered them from BRICKS. Reason: every module was discharged at a DEGENERATE WITNESS (witness collapse), not at any real $SU(3)$ operator with $T_{\text{real}} > 0$.

The three collapse shapes

SHAPE 1 -- Trivial satisfiability: 'EXISTS $C > 0$ AND $\mu > 0$ ' witnessed by $C = \mu = 1$. Says nothing about any correlator.

SHAPE 2 -- Reflexive bound: $\text{integrated_tail}(L, m) := \exp(-m \cdot L)$, so 'bound' $\exp(-m \cdot L) \leq \exp(-m \cdot L)$ closes by le_refl . No real heat-trace integral.

SHAPE 3 -- Unused hypothesis: mass gap hypothesis assumed, then $f(t) = \exp(-m \cdot t)$ discharged directly. Spectral hypothesis never used.

#	File	Named Open Prop	Collapse	$T_{\text{real}} > 0$
1	YM/ClusteringCore	clusters_zero_OPEN	const-zero pair	none
2	YM/MassGapStandin	massGap_standin_OPEN	zero CLM on C	none
3	YM/SpectralGapCore	hasMassGap_zero_OPEN	$T := 0$ (vacuous)	none
4	YM/TransferOperatorBound	transfer_gap_zero_OPEN	degen. operators	none
5	YM/TwoPointDecay	clustering_zero_OPEN	const-zero fn	none
6	YM/MassGapFromDecay	mass_gap_from_zero_OPEN	zero witness op	none
7	YM/IntegratedTailReal	integrated_tail_le_exp_OPEN	reflexive bound	none
8	YM/TransferGapReal	transfer_gap_real_OPEN	degen. witness	none
9	YM/MassGapReal	mass_gap_from_transfer_OPEN	zero on $H := C$	none
10	YM/ClusteringImpliesGap	clustering_implies_gap_OPEN	zero antecedent	none
11	YM/TransferImpliesClustering	transfer_implies_OPEN	const-zero fn	none
12	YM/TailImpliesTransfer	tail_implies_transfer_OPEN	definitional rewrap	none
13	YM/GapToDecay	gap_to_decay_OPEN	hypoth. unused	none

Property	Value
Modules audited	13
Non-vacuous $T_{\text{real}} > 0$ proofs	0
Converted to named open Prop	13
sorry / sorryAx / admit count	0 / 0 / 0
Axiom footprint (each)	{propext, Classical.choice, Quot.sound}
Compile (direct-lean, all modules)	EXIT 0
Surface #1 (Clay YM mass gap)	OPEN

12. The Conjecture Landscape -- Open Problems

The Wall256 project surfaces a rich landscape of genuine open conjectures. This section catalogues them with dependencies and what would be needed to resolve them -- from analytic (Varadhan) to number-theoretic (C13 law).

Conjecture C1: hw1 Bridge

STATEMENT: For all $\beta > 2.079416880124$: $w_1(\beta) < 1/7$ AS A LEAN PROOF TERM.

WHAT IS NEEDED: Requires SU(3) character theory or verified cubature in mathlib v4.12.0. Numerically certified (CERT_Arb). Currently absent from mathlib.

IF CLOSED: Would discharge H1. Still needs H2 (hOS) and H3 (h_bridge).

Conjecture C2: Osterwalder-Seiler Rate

STATEMENT: $\text{TruncatedActivityBound}(a)$ with $l > \log(7)$ STRICTLY, given $w_1 < 1/7$.

WHAT IS NEEDED: Requires formalizing Osterwalder-Seiler 1978 Thm 2.1 in Lean. The Ursell/truncated expansion is absent from mathlib. Hardest open leaf.

IF CLOSED: Would discharge H2. Still needs H3 and stays at lattice scope only.

Conjecture C3: Varadhan Small-t Asymptotic (Task #156)

*STATEMENT: $K(t) \leq C * t^{-4} * \exp(-c/t)$ for t in $(0, t_0]$ with $C, c, t_0 > 0$ explicit.*

WHAT IS NEEDED: Classical (Molchanov/Varadhan for compact semisimple Lie groups, exponent $-\dim(G)/2 = -4$ for SU(3)). No machine-checked proof exists in this project.

IF CLOSED: Closes gap between Surface #1 (done) and Surface #2 (blocked on this).

Conjecture C4: Per-Plaquette Activity Bound

*STATEMENT: EXISTS C, c, t_0 s.t. $\text{plaquette_activity_pw}(\beta, N, p) \leq C * \beta^{-4} * \exp(-c/\beta)$.*

WHAT IS NEEDED: Direct consumer of Varadhan (C3). Lean: exists_c_per_plaquette_pw (sorry).

IF CLOSED: Required for Kotecky-Preiss (C5).

Conjecture C5: Kotecky-Preiss Criterion at Small Coupling

*STATEMENT: EXISTS β_0 s.t. for $\beta \geq \beta_0$: KP condition holds with $\mu(\gamma) = a * |\gamma|$.*

WHAT IS NEEDED: Lean: kotecky_preiss_criterion (sorry, ClusterExpansion.lean line ~601). Needs polymer graph branching count and partial-sum estimate (neither formalized).

IF CLOSED: Would give convergence of the polymer log-partition series.

Conjecture C6: Area-Law Wilson Loop Decay

*STATEMENT: $|| \leq K * \exp(-m * \text{Area}(C))$ for contractible Wilson loops.*

WHAT IS NEEDED: Wilson loops are not defined anywhere in the codebase. Standard result from convergent cluster expansion at strong coupling (lattice only).

IF CLOSED: Lattice area law only. Continuum limit $a \rightarrow 0$ is a SEPARATE surface.

Conjecture C7: H4 / C13 Law (Requires Correction -- see Section 13)

STATEMENT: $\text{digit_len}(p) \geq 13$ implies $\text{Sym}(p) = 1$ FOR PRIME p .

WHAT IS NEEDED: Empirical: 6/6 cases. Named open Prop C13_Law_Open in H4_Derivation.lean. AUDIT-1 (Section 13): P5_genuine = 1000000001119 = 7 x 142857143017 is COMPOSITE. The law needs reformulation over primes only, and P5 needs a replacement.

IF CLOSED: After correction: law statement stays, but P5_genuine must be replaced.

Conjecture C8: Bessel I Formal Bound

STATEMENT: $I_n(x) \leq \exp(x) / \sqrt{2\pi x}$ for large x (standard asymptotic).

WHAT IS NEEDED: No Lean formalization in this project. Would provide a formal route to bounding the SU(3) character integrals and eventually to C1 (hw1 bridge).

IF CLOSED: Does not by itself prove $w_1 < 1/7$ as a Lean term; feeds toward C1.

Conjecture C9: Tighter Casimir Bounds (Task #157)

STATEMENT: $C_2(m,n) \geq c^(m+n)^2$ and $\dim(m,n) = O((m+n)^3)$ as Lean lemmas.*

WHAT IS NEEDED: Current formalization uses only linear/product-square slack versions. Tighter bounds are standard in the informal literature.

IF CLOSED: Tech-debt only. Would shorten Surface #1 proofs; no new theorem content.

13. Audit Findings (caught during pre-build certification)

The Opera Numerorum rule: errors are certified and documented, never silently corrected. The following two findings were discovered by running the pre-build certification script (certificates/certify_wall256_ym.py) against the uploaded source documents.

AUDIT-1: P5_genuine is COMPOSITE (not prime)

FINDING: P5_genuine = 1000000001119 = 7 x 142857143017. NOT PRIME.

TOWERS_YM_v2.3 states: 'P5_genuine = 1000000001119 is the ONLY live P5 in v2.3 -- the real 13-digit boundary prime (digit_len = 13, Sym = 1).'

CERTIFICATION RESULT (Miller-Rabin, witnesses [2,3,5,7,11,13,17,19,23,29,31,37]):

```
P5_genuine = 1000000001119
P5_genuine % 7 = 0
P5_genuine = 7 x 142857143017
Verification: 7 x 142857143017 = 1000000001119 [CONFIRMED]
Is 142857143017 prime? YES (Miller-Rabin)
CONCLUSION: P5_genuine = 7 x (prime) --> COMPOSITE.
```

The digit-count claim is still correct (13 digits, PASS). The primality claim is FALSE.

The Remarkable Structure of the Factorization

The factorization P5_genuine = 7 x 142857143017 has a striking feature:

```
1/7 = 0.142857142857142857... (the KP threshold denominator, repeating)

142857143017 begins with: 142857...
The cyclic group of 1/7 = {142857, 285714, 428571, 571428, 714285, 857142}

q = 142857143017 ~ 142857 x 10^6 + 143017
  = 142857 x (1000000 + 1.001...) -- very close to 142857 * 10^6

KP threshold: w1 < 1/7 = 0.142857...
P5_genuine = 7 x q where q ~ 10^11 * (1/7)
```

The 'boundary prime' is divisible by the very denominator that defines the KP convergence threshold. Whether this is a coincidence, a consequence of the construction of P5, or a deeper structural link to the C13 law -- this is an open question worth exploring (see Conjecture C7).

Required corrections

Document	Current claim	Corrected status
TOWERS_YM_v2.3	P5_genuine = 1000000001119 is a 13-digit boundary prime	P5_genuine = 7 x 142857143017 is COMPOSITE. A replacement prime with 13 digits is needed.
H4_Boundary.lean	C13 := 13 (boundary digit count, with P5_genuine as witness)	C13 is wrong; the PRIME witness for 13 digits needs replacement.
H4_Derivation.lean	C13_Law_Open: digit_len p >= 13 => Sym p = 1	Conjecture statement is fine; test case P5 needs a genuine prime witness.

AUDIT-2: beta_0 bracket requires the full 36-moment computation

FINDING: Our 7-moment partial sum (m_0..m_6) cannot reproduce the CERT_Arb bracket. The partial sum converges too slowly near beta_0.

The CERT_Arb computation (in the source workspace, TheoremaAureum143) uses N=36 terms with all exact moments m_0..m_36, computed by constant-term extraction over the SU(3) torus. The tail bound is 4.46e-32, giving sufficient precision to locate the 1/7 crossing to 12 decimal places.

Our 7-moment partial sum underestimates w1 near beta_0 because the missing terms m_7..m_36 each contribute positive amounts. The qualitative result (w1(0.86) >> 1/7, D4 clearly fails) is confirmed even with 7 terms. The precise bracket requires the full computation.

Computation	Moments used	Precision near beta_0	D4 status	Verdict
-------------	--------------	-----------------------	-----------	---------

Our script (7 terms)	m_0..m_6 (exact)	Insufficient (underestimates w1) w1(0.86) >> 1/7: CONFIRMED	Bracket: NEEDS FULL CERT_ARB
CERT_Arb (authoritative)	m_0..m_36 (exact rational)	fail <= 4.46e-32 (dps=80)	w1(0.86)=0.432367... > 1/7 CERTIFIED: beta_0 in [2.079416880123,

CERT_Arb SHA: b5a9f0a7666a91f283a7d4531ae99dff2097c2cedef10424b77833b7bbc840d3

D4 failure (w1(0.86) >> 1/7) confirmed by our independent 7-term check.

14. What IS and IS NOT Proven -- Clay Rigour Checklist

WHAT IS PROVEN (machine-checked, trio-clean)

kp_summable_of_truncatedActivity: IF TruncatedActivityBound(a) AND $N(n) \leq 7^n$, THEN the entropy-weighted polymer series is summable. [Wall256_Note.lean]

mass_gap_pos_of_spectral_gap: $\rho^d = \exp(-\Delta \cdot d)$ algebra. [Wall256_Note.lean]

strong_coupling_decay_of_open_inputs: threads H1+H2+H3 into abstract lattice decay. SORRY=0.
Axioms={propext,Classical.choice,Quot.sound}. [Wall256_Scaffold.lean]

PeterWeyl_Summable_SU3: $\sum_{(m,n)} \dim(m,n)^2 \exp(-t \cdot C2(m,n))$ converges for $t > 0$.

Truncation-to-envelope bridge $S_N(t) \leq K(t)$. [Tasks #154, #155]

beta_0 in [2.079416880123, 2.079416880124]: CERTIFIED (CERT_Arb, out-of-tower).

D4 at beta=0.86: w1(0.86) >> 1/7 CERTIFIED. D4 FAILS.

All 20 Bessel $I_n(x)$ values verified against Z_BESSEL_I_COMPLETE CSV: ALL PASS.

Monotonicity of w1(beta): strictly decreasing (Laplace transform argument).

Strict-inequality threshold: $l = \log(7)$ is the exact divergence boundary.

$S_N(t=1)$ convergence: series stabilises at N=5 to 6 decimal places.

P5_genuine digit count: 13 digits -- PASS.

WHAT IS NOT PROVEN -- explicit Clay-scope list

NO Yang-Mills mass gap. No $\mu > 0$ unconditionally. $\Delta > 0$ exists only under H1+H2+H3.

NO spectral gap for any real Wilson transfer operator. corr/sep are abstract; no real SU(3) operator built.

NO closure of Clay Surface #1. YM tower Status: OPEN.

NO continuum limit statement. Scope: LATTICE SU(3) strong-coupling only.

NO convergence of any real cluster expansion. KP conditional on H2's rate.

NO discharge of H1 (hw1), H2 (hOS), H3 (h_bridge). None scheduled.

NO cross-tower bridge between YM and NS reductions.

NO RH / GRH claims in YM certs. M7/GRH/C05 tower: SEPARATE.

P5_genuine: NOT PRIME. (AUDIT-1: = 7 x 142857143017.)

15. Bessel I Zero-Pattern: LLM Error Topology on the Z-Protocol

This section reports an empirical study of how a language model (LLM) reproduces modified Bessel function values $I_n(x)$ under the Z-protocol (T=0: LLM-direct, no computational tools; T=1: tool-assisted). The zero-error cells -- where the LLM gets the value exactly right -- are not randomly distributed. They trace out a specific locus determined by the Sym parameter and by the ratio x/n .

Experimental setup

Parameter	Values tested	Trials per cell
Order n	0, 1, 2, 3, 4, 5	5 (T=0 LLM-direct)
Argument x	0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 5.0, 10.0	5
Sym	1 (n=0,1); 2 (n=2,3,4,5)	--
Method T=0	A_LLM_T0: LLM answers with no tools	--
Method T=1	B_tool_T1: mpmath / rule echo	100 (all ZERO)
Source	Bessell_Z_MEASURE_(2) CSV	--

Zero map (T=0 LLM-direct): ZERO = correct / ERR = wrong / prt = partial

Function	x=0.5	x=1.0	x=1.5	x=2.0	x=2.5	x=3.0	x=5.0	x=10.0
I_0 (Sym=1)	ZERO	ZERO	ZERO	ZERO	prt	ZERO	ZERO	ZERO
I_1 (Sym=1)	ERR	ERR	ZERO	ERR	ERR	ZERO	ERR	ERR
I_2 (Sym=2)	ERR	ZERO	ERR	ERR	ERR	ERR	ERR	ERR
I_3 (Sym=2)	ERR	ERR	ERR	ERR	ERR	ERR	ERR	ERR
I_4 (Sym=2)	ERR	ERR	ERR	ERR	ERR	ERR	ERR	ERR
I_5 (Sym=2)	ERR	ERR	ERR	ERR	ERR	ERR	ERR	ZERO

Aggregate zero-error rate by Sym

Sym	n values	Total cells	LLM zeros (correct)	Full errors	Partial
1	n=0, n=1	16	9 (56.2%)	6 (37.5%)	1 (6.2%)
2	n=2..n=5	32	2 (6.2%)	30 (93.8%)	0 (0.0%)

The 11 zero-error cells and their x/n ratios

Cell	Sym	$I_n(x)$ value (15 dp)	x/n ratio	Ratio type
I_0(0.5)	1	1.063483370667600	infy (n=0)	n=0: ratio undefined
I_0(1.0)	1	1.266065877752008	infy	n=0: all x rational
I_0(1.5)	1	1.646723189621999	infy	n=0: all x rational
I_0(2.0)	1	2.279585302336067	infy	n=0: all x rational
I_0(3.0)	1	4.880792585865024	infy	n=0: all x rational
I_0(5.0)	1	27.239871823604449	infy	n=0: all x rational
I_0(10.0)	1	2815.716628466254	infy	n=0: all x rational
I_1(1.5)	1	0.981666428761618	1.5 = 3/2	half-integer
I_1(3.0)	1	3.953370217402609	3.0 = 3/1	integer
I_2(1.0)	2	0.135747669767038	0.5 = 1/2	half-integer
I_5(10.0)	2	777.188286403312	2.0 = 2/1	integer

Key findings

T=1 (tool-assisted): 100% zero-error across ALL 48 cells. Tool use eliminates Bessel I reproduction errors entirely.

T=0, Sym=1: 56.2% zero-error rate. The LLM has internal representation of I_0 and I_1 across most of the tested range.

T=0, Sym=2: 6.2% zero-error rate. The LLM is almost entirely blind to I_n for n >= 2, except at two cells.

The two Sym=2 zero-error cells (I_2(1.0), x/n=0.5) and (I_5(10.0), x/n=2.0) both fall at integer or half-integer x/n ratios.

All non-n=0 zero-error cells (I_1, I_2, I_5) have x/n in {1/2, 3/2, 2, 3} -- rational numbers with small denominator.

The single partial-error cell I_0(2.5): x=2.5 is not a round number. 3 errors out of 5 (60% error rate) even for the best-known order n=0.

Interpretation

The LLM's zero-error locus is NOT random. It follows two rules:

RULE 1 (Sym): Sym=1 inputs (n=0,1) yield far more zeros than Sym=2 inputs. The model's internal Bessel knowledge is concentrated at low order.

RULE 2 (Ratio): Among higher-order functions (n >= 1), zeros appear only where x/n is a rational number with denominator 1 or 2. The model fails at irrational or non-simple-rational x/n points.

COMBINED: The LLM error function on Bessel I space has a zero-set that concentrates on Sym=1 AND rational x/n. This is a knowledge-topology map, not a random noise floor. An LLM operating without tools is reliable on I_0 (the most-tabulated order) but structurally unreliable on I_n (n >= 2) except at the simplest rational-ratio arguments.

Connection to the Wall256 / YM project

The Bessel I barrier (Open Question OQ-BI) requires formal I_n(x) bounds in Lean / Mathlib. The Z-protocol finding reinforces why: an LLM-only proof attempt on Bessel I would have a 93.8% error rate for Sym=2 inputs. The barrier is real. Tool-assisted computation (T=1) gives 0% error. The correct route is machine verification, not LLM recall.

Source: Bessell_Z_MEASURE_(2)_1780557621561.csv. All T=1 verification against Z_BESSEL_I_COMPLETE CSV: ALL 20 values PASS (relative error < 1e-10).

16. Cryptographic Binding and Sources

Artifact	SHA-256 / Source
CERT_Arb_beta0 (interval certificate)	b5a9f0a7666a91f283a7d4531ae99dff2097c2cedef10424b77833b7bbc840d3
D4_w1_NEGATIVE_Certificate	9a794ccf0c707812e6fa3db2095a350f2d5b61a011fcc77e453c548716ac8764
D1_to_D3_Plan.md	a40422fe7595ebc2c3c7d8d4b80a2c5645431ac1a52f3699c732562d4574b553
Big_Day_Status_2026-06-01.md	9a00f1106bf93da155d8366521579eeec17bcacd2f4aea62fbcd0159dedfc7d1
hits.txt Genesis seal (460 bricks)	eeebcd9a540aa7a2c90edd23827c73e4d1bb5af641d352f70a5de849b21f875f
Lean commit (Wall256 scaffold)	8eeab54
Lean commit (Beta0Certified wiring)	ac241de

Per-file SHA-256 of the 13 YM/OS modules (from PDF 263, 2026-06-03)

#	File	SHA-256
1	YM/ClusteringCore.lean	f3ceecd3892edd81a760aa8205e947aef57ee273c84ef5be33c3300f8aa9a1fb
2	YM/MassGapStandin.lean	ffa6f5820e6dfcd6b0af4051582be2508a35e21117d87e24f26eb799e8ee5d
3	YM/SpectralGapCore.lean	fca1de2b085e768a61ef320db8d7d23f546a0f3ddce0b76ccbc3f956eee69e89
4	YM/TransferOperatorBound.lean	ff0bb4887885a10e11ef0e41a46dd801acb0ff5cb58b75d3d613a3d27e10295c
5	YM/TwoPointDecay.lean	516b8fb721a762181a0200f11e8fd01f16ee01c19ae8228fca944ab907234575

6	YM/MassGapFromDecay.lean	b5acbbc0770f6211765989440ec82a88d6d2744ae7a5d45d383eaa1360adaa3b
7	YM/IntegratedTailReal.lean	223a0dacab5f7bc4b871b9a1dbe620d7944af091d7bbd65d6f6f693d68ef8ca2
8	YM/TransferGapReal.lean	bd30cd0f17d3a74ae8870c1fcdf2808a9584d6e77185274760367abda5484ebe
9	YM/MassGapReal.lean	30f0b1b37c510cee3265a1f5e3e2bccbf04647a500273f28a09902ffc51fb0
10	YM/ClusteringImpliesGap.lean	acc573093f6dc6b68e7496ebfd99932ba91a63937a5430c5f4e5eadea3ba1626
11	YM/TransferImpliesClustering.lean	3e502bacd4110ed8e52b5ad1841e98c0e45164e974756b013957b092bff7a301
12	YM/TailImpliesTransfer.lean	e0d3ba5d3cbcad9f17cea42a73cc155af00af2dbb91fc0254f70c8016643bd50
13	YM/GapToDecay.lean	ff4b8f5dd357a353982afe08e66afd8036d4eabc526af4d0c8c364828f235dbe

Generated 2026-06-04 | Opera Numerorum | David J. Fox | ORCID 0009-0008-1290-6105

YM_STATUS = OPEN | SORRY = 0 | COMMIT = 8eeab54 | Mathlib v4.12.0

No mass-gap / $\mu > 0$ / Surface-#1-closed / Clay / RH claim is made.

Two audit findings documented above. Errors are certified, not hidden.